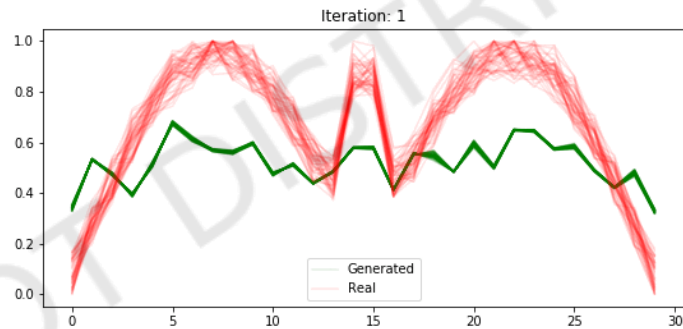


Generative adversarial network

Generative adversarial networks (GAN) are models used to **generate sample data of interest** (rather than discriminative purposes such as classification / scoring).



Generative adversarial network

It is based on the idea* that to **training generative models using discriminatory function as supervisor**.

The system has two actors:

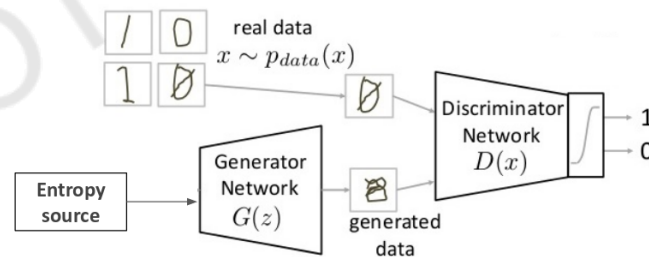
- A "*generator*" function
 - with trainable parameters to generate sample data mimicry to training data
- A "*discriminator*" function
 - capable of measuring the difference between generated data and training data
 - usually also with trainable parameters

(* special note: The idea is neither novel nor restricted to neural networks.)

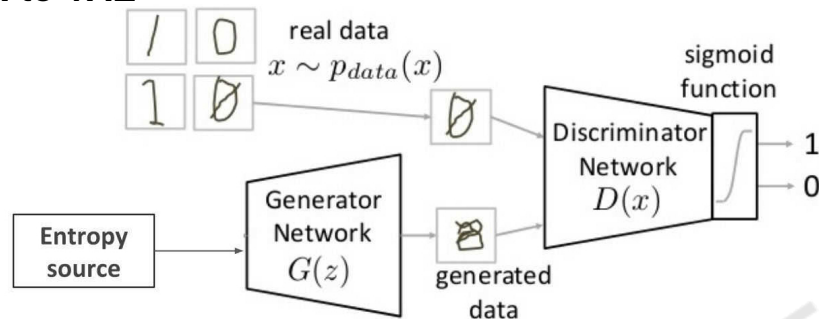
Generative adversarial network

Steps

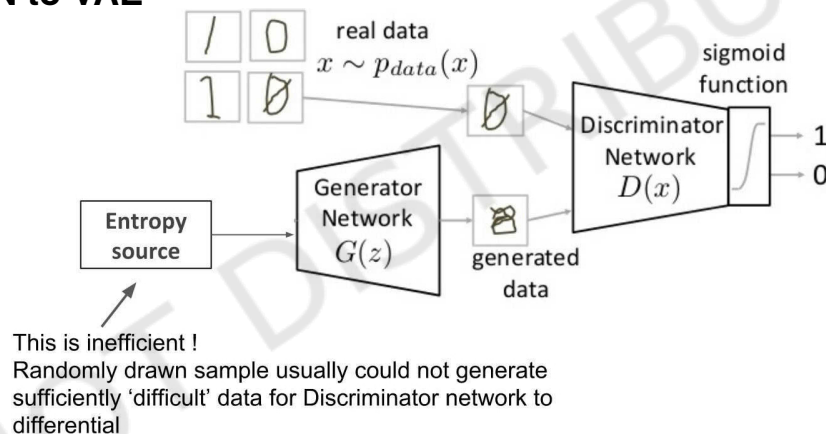
- *Generator* draws (random) inputs from an "entropy source" (e.g. a random number generator) and generates output samples from it
- *Discriminator* computes an error measure (e.g. binary cross entropy or softmax loss) for the each generated output sample against real data
- The errors (dissimilarity between generated and real data) are back-propagated through the *generator* to fit the parameters.
- (The generated output samples are used, together with the real data, to improve the *discriminator*)
- Repeat until the error is minimized.



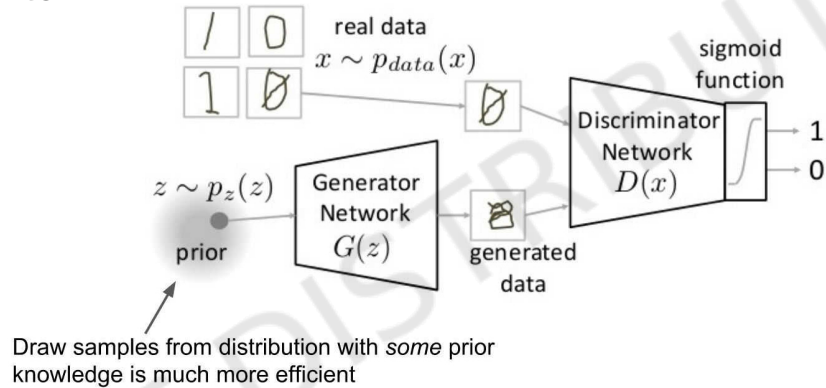
From GAN to VAE



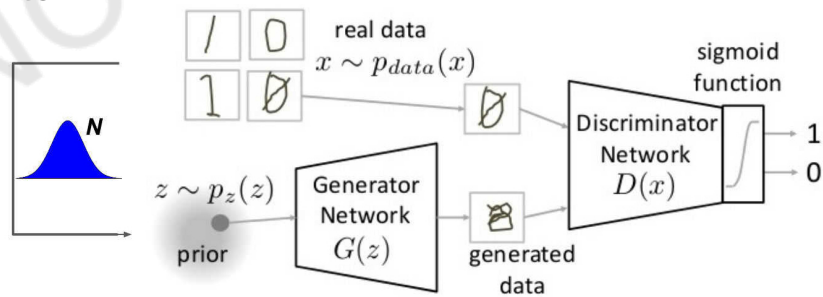
From GAN to VAE



From GAN to VAE



From GAN to VAE



Prior distribution can be learned from training set data

- Prior(s) should be easily computable, e.g. Gaussian form

Variational (Bayesian) Inference formulation

Sampling X over latent variable Z

$$P(X) = \int P(X|z)P(z)dz$$

Use distribution Q to approximate P

$$Q(z|X) \sim P(z|X)$$

"Difference" between P and Q can be measured with K.L. divergence:

$$\text{KL}(Q(z|X)||P(z|X)) = E_{z \sim Q} [\log Q(z|X) - \log P(z|X)]$$

(applying Bayes' rule)

$$= E_{z \sim Q} [\log(Q|X) - \log P(X|z) - \log P(z)] + \log P(X)$$

Rearranging gives

$$\log P(X) - \text{KL}(Q(z|X)||P(z|X)) = E_{z \sim Q} [\log(Q|X) - \log P(X|z) - \log P(z)]$$

$$= E_{z \sim Q} [\log P(X|z)] - \text{KL}(Q(z|X)||P(z))$$

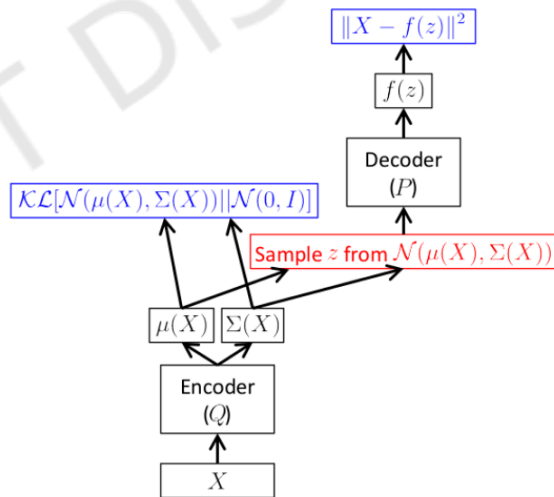
Analytic solution for Gaussian case

if $P(z) \sim N(0, I)$ and $Q(z|X) \sim N(z; \mu, \sigma^2)$

KL has an analytic solution:

$$\text{KL}(Q(z|X) || P(z)) = -\frac{1}{2} \sum_{j=1}^J (1 + \log \sigma_j^2 - \mu_j^2 - \sigma_j^2)$$

Variational Autoencoder



Paradigm shift on universal function approximator

Previously

- Any distributions could be approximated with neural network

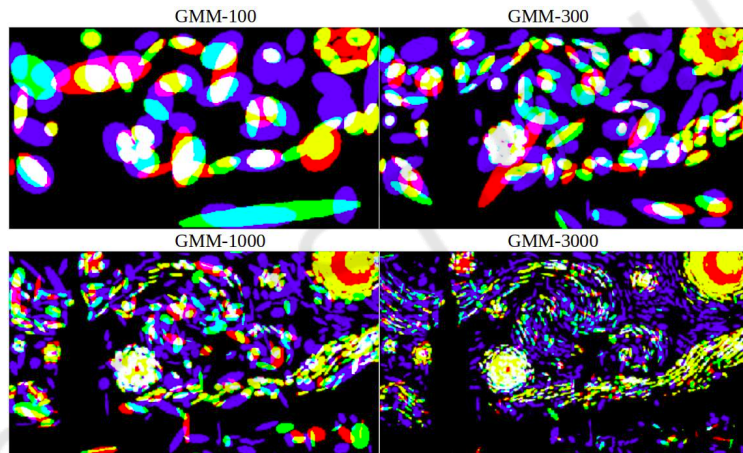
VAE

- Any distributions could be approximated with **high-dimensional Gaussian distribution**
- Parameters of the Gaussian distribution can be learned by neural network

Future (?)

- What other distributions can we use ?

High dimensional Gaussian approximator

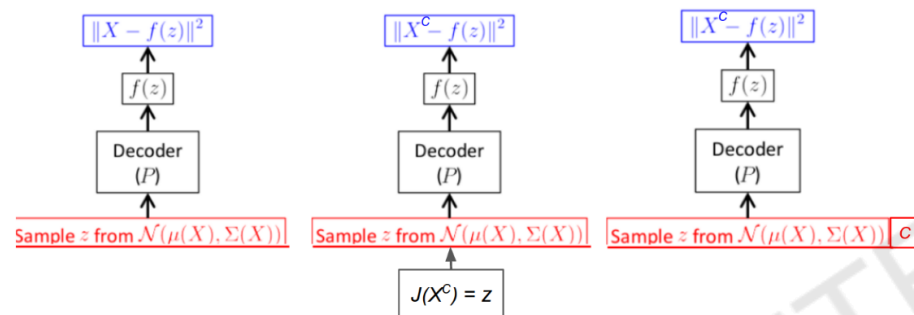


Conditional Variational Autoencoder (CVAE)

- In the vanilla implementation of VAE, one cannot control the output

Easy solution

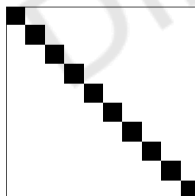
- Recognize that each output corresponds to an input
 - (drawn from an entropy source (e.g. a random number generator))
- Train an encapsulating network to relate between input and output label



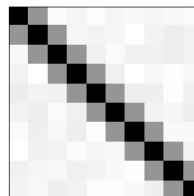
(extra) Recurrent latent Gaussian model

- Remember we draw z from $P(z) \sim N(0, I)$
- Assume each latent Gaussian component is independent to one another
- The constraint was introduced for computational purpose
- N -dim Gaussian: N means, $N(N + 1)/2$ covariances to model
 - at dim=1000, full covariance matrix has 501,500 parameters; diagonal matrix only with 2000 parameters
- Can we model a bit more parameters (in between diagonal and full) ?

Diagonal



Diagonal+1



- Make a Gaussian component dependent on its previous component
 - Essentially a *recurrent Gaussian model*