

Research trend - Algorithm Cocktails



Monte Carlo Tree Search (MCTS)

- An exhaustive "full" tree search would solve the decision problem

Example: Tic-Tac-Toe

- Maximum of 9 steps
- Maximum of 9 nodes


Pseudo-code for full search tree on tic-tac-toe

```
for step in [1..9]:
    new_layer = List()
    player *= -1
    for node in last_layers:
        for i in [1..9]:
            state = node.state
            state[i].set_color(player)

            if is_valid(state) {
                _node = new node(state)
                node.add_child(_node)
                new_layer.add(_node)
            }
        }
    }
    layers.add(new_layer)
}
```

- Generates a tree of 9 layers with 549,946 leaf nodes

Tic-Tac-Toc full tree



Monte Carlo Tree Search (MCTS)

Core idea:

- An exhaustive "full" tree search would solve the decision problem
 - (except that it is computationally prohibitive)
- Expand tree only up to a few layers (steps)
- Identify only the more **valuable** nodes to sample, criteria being
 - [exploitation] known to give higher estimated reward, or
 - [exploration] higher uncertainty of the upper reward bound
- Given the resources (e.g. time), sample and refine the estimated value of each move
- Finally, choose the move with the highest refined estimated value

Monte Carlo Tree Search (MCTS)

Core idea:

- An exhaustive "full" tree search would solve the decision problem
 - (except that it is computationally prohibitive)
- Expand tree only up to a few layers (steps)
- Identify only the more **valuable** nodes to sample, criteria being
 - [exploitation] known to give higher estimated reward, or
 - [exploration] higher uncertainty of the upper reward bound
- Given the resources (e.g. time), sample and refine the estimated value of each move
- Finally, choose the move with the highest refined estimated value

Components

1. Estimator on the reward
2. Estimator on the upper reward uncertainty bound
3. A way to sample and refine the estimated value

Monte Carlo Tree Search (MCTS)

Estimator of reward

The reward estimator updates as simulation proceeds, and usually takes the form of

$$Q(s, a) = \frac{w(s, a)}{n(s, a)}$$

where s is current state, a is the move action, $w(s, a)$ is the winning count and $n(s, a)$ is the observation count.

Estimator of uncertainty bound

- This purpose of this term is to encourage sampling on less visited nodes.
- It should scales with total number of samples N , and inversely with current node visit count n_i , e.g.

$$\mu(s, a) = c \frac{N^\alpha}{n^{\beta(s, a)}}; N = \sum_b n(s, b)$$
- A commonly used expression is $\mu(s, a) = c \sqrt{\frac{\log N}{n(s, a)}}$
- it is possible to use improve the estimator by assigning a *prior* on this bound if *we have some knowledge on the context* (s, a).

A way to sample and refine the estimated value

- Usually the expression of uncertainty bound is unchanged $\mu(s, a)^{\text{new}} = \mu(s, a)^{\text{old}}$
- Estimation of reward is updated with a **simulation**

- simulate the final outcome (win/loss) after choosing the move

$$Q^{\text{new}}(s, a) = \frac{w(s, a) + v}{n(s, a) + 1}; v = \begin{cases} 1 & \text{if win} \\ 0 & \text{if loss} \end{cases}$$

- In case without any prior knowledge, a "rollout" (random playout of the game) can be used for the simulation

Monte Carlo Tree Search (MCTS)

Vanilla implementation

1. At every iteration we compute UCT for every move a of the current state s as

$$U(s, a) = Q(s, a) + \mu(s, a) = \frac{w(s, a)}{n(s, a)} + \mu(s, a)$$

2. We then **select** the move a_t that maximizes the $U(s, a)$ and transverse to the next (leaf) node s_L .

$$a_t = \underset{a}{\operatorname{argmax}} U(s, a)$$

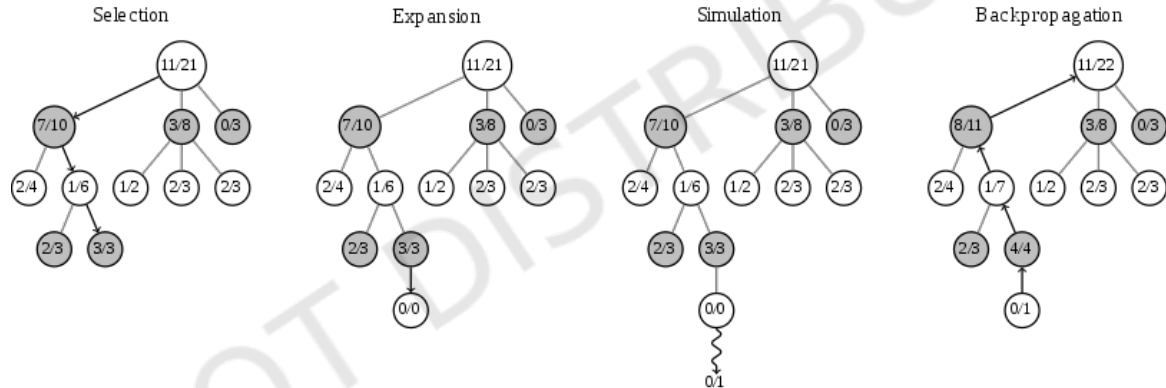
3. (We may optionally **expand** the leaf node. The number of level expansion is called the "depth" of the tree)
4. Start a **simulation** from s_L until end state
5. $Q(s, a)$ (and hence $U(s, a)$) value is updated by **backpropagation** through the node trail based on the simulation outcome v

$$Q^{\text{new}}(s, a) = \frac{w(s, a) + v}{n(s, a) + 1}; v = \begin{cases} 1 & \text{if win} \\ 0 & \text{if loss} \end{cases}$$

6. Repeat the process as resources allow

Monte Carlo Tree Search (MCTS)

Schematic diagram for MCTS



AlphaZero

There are two opportunities for function approximator in MCTS algorithm:

1. Estimator of uncertainty bound $\mu(s, a)$
2. Simulation $v(s)$

Estimator of uncertainty bound

- AlphaZero uses the expression:

$$\mu(s, a) = \frac{c\sqrt{N}}{1 + n(s, a)} \cdot \boxed{P(s, a)}$$

- The first term encourages sampling on less visited nodes
- The second term is a *prior* modifier of the action selection

A way to sample and refine the estimated value

- In AlphaZero, the final outcome is directly predicted from a neural network
 - No explicit simulation is performed
 - (explicit simulation was performed in AlphaGo implementation)

$$Q^{\text{new}}(s, a) = \frac{w(s, a) + \boxed{v(s)}}{n(s, a) + 1}$$

The boxed terms are obtained from a neural network

AlphaZero

AlphaGo networks

- In AlphaGo implementation, two neural networks were built:
 - $P(s, a)$ is obtained from a "policy network"
 - $v(s)$ is obtained from a "value network"
- Policy network $\text{net}_{\text{policy}}$ takes the current game state ($19 \times 19 = 361$ positions with 3 possible values {black, white, empty}) as input, and computes a 362 (=361 positions + pass) real-valued vector as output.
- Value network $\text{net}_{\text{value}}$ takes the same input and computes a real-value vector as output.

$$S = \{0, 1, -1\}$$

$$\text{net}_{\text{policy}} : S^{361} \rightarrow R^{362}$$

$$\text{net}_{\text{value}} : S^{361} \rightarrow R$$

AlphaZero network

- In AlphaZero implementation, **the two networks are merged into one.**
- The network takes the same input, but output a 363 real-valued vector (361 positions + pass + value output)

$$\text{net}_{\text{zero}} : S^{361} \rightarrow R^{363}$$

- The loss function l combines mean-squared error from value network and cross-entropy losses from policy network

$$l = (z - v)^2 - \pi^T \log P + c ||\theta||^2$$

where (z, π) are the actual outcome and probability of moves from training data respectively

(* in actual implementation, 7 steps are used in order to resolve the three-way knot scenario in Go)

AlphaZero

Reinforcement Learning

- MCTS algorithm could generate gameplays by self-play
- Gameplays are then utilized as training data

AlphaZero

Other details

- Artificially encourage more explorations
 - Stochastic policy with temperature control
 - Dirichlet prior
- Symmetry
 - Rotational and reflection invariance are utilized
- Neural network model succession
 - A new network is pitched against its successor
 - New network takes over if winning rate > threshold (55% in AlphaZero)
- Asynchronous
 - self-play, neural network training and model succession are performed in parallel
- Hardware / Distributed computing
- ...